

Physics-Closed Multimodal Dense Prediction via Latent Continuum Equivalency: Factor-of-Safety Cells and Mechanistic Path Equilibrium Fusion

Samreedh Bhuyan^{1*} Dr. Anil Earnest²

¹School of Computer Engineering, Kalinga Institute of Industrial Technology (KIIT), India

²CSIR Fourth Paradigm Institute (CSIR-4PI), Bengaluru, India

July 31, 2026




Figure 1: Opening illustration of dense multimodal segmentation under anisotropic optical-topographic fusion (RGB, ground truth, prediction, confidence heatmap). Landslide inventory mapping is used as a stress test of the proposed physics-closed continuum operators.

Abstract

Multimodal dense predictors often fuse heterogeneous streams with unconstrained statistical gates, even when one modality encodes a continuum mechanical field rather than an appearance cue. We ask whether a single reduced-order continuum operator can be injected into high-dimensional networks without destabilizing backpropagation, and whether that operator can serve as a commensurable currency for both feature gating and dual-path fusion. We introduce **Latent Continuum Equivalency (LCE)**: the same Taylor-stabilized infinite-slope factor-of-safety (FS) skeleton appears at the encoder inlet, hybrid bridges, skip lattice, dual physics decoders, and output routing. Material parameters use bounded exponentials $c = \exp(\text{clip}(w))$, which we prove place resisting/driving forces on a compact set and yield a locally Lipschitz FS gate. Dual-path logits are merged by **Mechanistic Path Equilibrium Fusion (MPEF)**, the softmax of failure energies, which we identify as the unique maximizer of an entropy-regularized linear energy on the simplex. Instantiated as GeoPhysicsLandslideNet (PS-GPLNet)—a pentastream CNN-physic-foundation architecture—the method is evaluated on Bijie and Landslide4Sense, anisotropic topographic-spectral segmentation benchmarks with severe imbalance. PS-GPLNet attains best-epoch F1/IoU of 0.933/0.907 on Bijie and 0.736/0.660 on Landslide4Sense, exceeding strong CNN, transformer, and dual-stream baselines under a unified protocol. Empirical GPU probes confirm material compactness, local FS-gate sensitivity bounds, and the MPEF energy identity. Inference studies further show structured fuse3 latents (t-SNE), Boundary F-score/Hausdorff fidelity, and flatter F1 decay under RGB+DEM Gaussian noise than U-Net and DeepLabV3+.

1 Introduction

Dense prediction from multiple sensors is a core problem in computer vision [1, 2, 3]. When modalities are correlated—as in RGB-D—learned fusion can exploit statistical redundancy. When modalities are *anisotropic*, one stream may encode an appearance field while another encodes a physical continuum (slope, shear, height), and unconstrained gates are free to ignore mechanics that the task actually depends on [4, 5].

A growing line of physics-informed learning places continuum or geometric constraints in losses or attention [6, 7, 4, 8, 5, 9]. Most susceptibility-oriented geo-AI work still keeps physics at the objective [10, 11], while dual-stream landslide segmenters keep fusion statistical [12, 13]. We instead close a reduced-order continuum loop *inside* the forward pass.

Our contribution is a general operator family for multimodal dense prediction. We formalize Latent Continuum Equivalency (LCE): every stage applies an operator from the same infinite-slope FS class, so failure energy remains a shared physical currency from inlet to outlet. Bounded material reparametrization keeps FS finite and the sigmoid gate locally Lipschitz. MPEF routes dual decode paths by softmax failure energy and is the unique entropy-regularized energy maximizer on the simplex [14]. Landslide segmentation on Bijie [15] and Landslide4Sense [16] is the stress test: spectral ambiguity plus topography governed by gravity/shear, not RGB-D correlation. Under a unified protocol, PS-GPLNet reaches F1/IoU 0.933/0.907 (Bijie) and 0.736/0.660 (L4S).

2 Related Work

Dense segmentation and multimodal fusion. U-Net [1], DeepLabV3+ [2], SegFormer [3], and TransUNet [17] dominate dense masks. Multimodal optical-DEM fusion improves landslide mapping [15, 16, 13, 18]. Dual-stream gated U-Nets [12] learn scalar gates; we use them only as empirical baselines. Interpretable fusion via unfolded sparse coding [9] motivates treating fusion as an optimization identity rather than a free MLP.

Geometric and topological priors in vision. Virtual Normal enforces high-order 3D geometry beyond pixel errors [4]. Topological PH losses inject discrete structure with differentiability arguments [8]. Lipschitz certificates for segmentation quantify robustness [19, 20]. Our FS gate plays an analogous role for continuum stability.

Physics inside architectures. Gravityformer multiplies transformer attention by a closed-form gravity interaction [5]. PINNs place PDE residuals in losses [6, 7]. Landslide PINNs largely target susceptibility [10, 11, 21]. PS-GPLNet differs by architectural closure of FS cells across encoding, fusion, decoding, and path equilibrium, with a foundation stream (Prithvi+LoRA [22, 23]).

Table 1: Conceptual positioning for dense multimodal prediction.

Approach	Dense mask	Continuum in forward	Foundation
CNN / ViT segmenters	✓	–	optional
Dual-stream gated [12]	✓	–	rare
Susceptibility PINN / PGML	often raster	loss / head	rare
Gravity-/geometry-informed CV [5, 4]	task-dep.	attention / loss	rare
PS-GPLNet (ours)	✓	FS cells + MPEF	Prithvi+LoRA

Positioning. Table 1 contrasts dense masks, FS-in-forward, and foundation context. The novelty claim is LCE + MPEF, not a new remote-sensing dataset.

3 Method

We present a physics-closed multimodal dense predictor whose forward pass is organized by a single reduced-order continuum narrative: the infinite-slope factor of safety (FS). The classical geotechnical FS is *not* wired verbatim into neurons. We first recall the textbook form, then derive a Taylor-stabilized neural operator FS_{nn} that relocates pore pressure into an additive driving term and admits a latent variant without trigonometric angles. That operator family is then realized end-to-end in GeoPhysicsLandslideNet (PS-GPLNet) as four stages:

- **Stage A (inlet):** five parallel encoders + physics proxies + pixel/latent mechanistic cells;
- **Stage B (hybridization):** Complementary Modality Bridges (CMB) that resonate CNN texture with physics pyramids;
- **Stage C (tri-stream fusion):** MAO-GeoEGCA at the bottleneck and TTEB multi-scale skips;
- **Stage D (dual decode + equilibrium):** dual PhysicsDecoders with PGDI, then MPEF instability routing at three supervision heads.

The remainder of this section follows that order after the FS derivation.

3.1 Problem and notation

Let $x_1 \in \mathbb{R}^{B \times 3 \times H \times W}$ denote the optical stream (RGB) and $x_2 \in \mathbb{R}^{B \times 3 \times H \times W}$ the topographic/auxiliary stream. On Landslide4Sense, x_2 stacks NDVI, slope, and DEM; on Bijie, DEM is replicated across three channels. The prediction target is a binary landslide mask $y \in \{0, 1\}^{B \times 1 \times H \times W}$. Unless stated otherwise we use $H = W = 256$ at the network inlet and a unified fusion width $C = 64$.

From (x_1, x_2) the network extracts normalized physics proxies (s, d, v) —slope, elevation, and vegetation—and maps them to geotechnical variables (α, h, m) :

- α : slope angle in radians (pixel-wise);
- $h > 0$: soil-column / elevation height proxy;
- $m \in (0, 1)$: moisture-saturation proxy.

Separate RGB-path and DEM-path mappers produce stream-specific triples (α_r, h_r, m_r) and (α_d, h_d, m_d) . We write $\odot$ for Hadamard product, $[\cdot \parallel \cdot]$ for channel concatenation, and $\varepsilon > 0$ for a numerical floor (typically 10^{-6}).

3.2 Classical infinite-slope factor of safety

For an infinite slope with seepage parallel to the surface, the textbook factor of safety compares resisting and driving forces on a potential slip plane of depth H [24]:

$$\text{FS}_{\text{class}} = \frac{c' + (\gamma_{\text{sat}}H - u) \cos^2 \alpha \tan \phi'}{\gamma_{\text{sat}}H \sin \alpha \cos \alpha}, \quad (1)$$

$$u = m \gamma_w H \cos^2 \alpha.$$

Here c' is effective cohesion, ϕ' the effective friction angle, γ_{sat} the saturated unit weight, γ_w the unit weight of water, α the slope angle, and $m \in [0, 1]$ the saturation ratio. Substituting u yields

$$\text{FS}_{\text{class}} = \frac{c' + (\gamma_{\text{sat}} - m\gamma_w)H \cos^2 \alpha \tan \phi'}{\gamma_{\text{sat}}H \sin \alpha \cos \alpha}. \quad (2)$$

Failure is declared when $\text{FS}_{\text{class}} \leq 1$. Equation (2) is a valid engineering calculator for hand analysis, but it is a poor neuron. The term $-m\gamma_w$ subtracts from the resisting numerator: early in training, uncalibrated material weights can drive that numerator through zero or negative values, flipping gradient signs and producing NaNs. Moreover, embedding $\sin \alpha$ and $\cos \alpha$ into deep latent tensors—where channels no longer represent geodetic angles—is physically meaningless and numerically brittle. We therefore derive a neural form that (i) relocates hydrology into an additive driving term and (ii) admits a latent counterpart that drops angles entirely while preserving the resisting/driving ratio skeleton.

3.3 Derivation of the Taylor-stabilized neural FS

A dense predictor scores *instability*. Define stress as the reciprocal $S = 1/\text{FS}_{\text{class}}$. Expanding (2),

$$S = \frac{A}{B - C}, \quad (3)$$

with the grouped forces

$$\begin{aligned} A &= \gamma_{\text{sat}}H \sin \alpha \cos \alpha, \\ B &= c' + \gamma_{\text{sat}}H \cos^2 \alpha \tan \phi', \\ C &= m\gamma_w H \cos^2 \alpha \tan \phi'. \end{aligned} \quad (4)$$

Here A is gravity drive, B dry structural resistance, and C hydrological destabilization. Whenever $C \geq B$, the denominator of (3) vanishes or turns negative—catastrophic for backpropagation.

Proposition 1 (First-order Taylor relocation of pore pressure). *Assume $C/B < 1$ in a neighbourhood of dry soil ($C = 0$). Then*

$$\begin{aligned} \frac{A}{B - C} &= \frac{A}{B} \frac{1}{1 - C/B} \\ &= \frac{A}{B} \sum_{k=0}^{\infty} \left(\frac{C}{B}\right)^k \\ &= \frac{A}{B} + \frac{AC}{B^2} + O((C/B)^2). \end{aligned} \quad (5)$$

Truncating after the linear term yields the first-order structural equivalence

$$\frac{A}{B - C} \propto \frac{A + C}{B} \quad (6)$$

up to a positive local scale B^{-1} . Consequently, subtracting hydrology from resistance is—to first order—equivalent to adding hydrology into the driving reference frame.

Proof. Factor B from the denominator of (3) and apply the geometric series $1/(1 - x) = \sum_{k \geq 0} x^k$ for $|x| < 1$ with $x = C/B$. Truncation after $k = 1$ gives (5). Multiplying numerator and denominator of the truncated form by B rearranges to the proportional additive drive (6). $\square$

Physically, the classical layout treats water as a lubricant that reduces effective normal stress; the neural layout treats water as an additive downslope driver. Both describe the same tipping point—driving forces overtaking resistance—but the neural layout keeps all tensors strictly positive. Reverting from stress S to an FS-shaped ratio and absorbing material constants into learnable positive weights (c, ϕ, γ, w_m) produces the operator used in pixel cells and in MPEF.

Definition 1 (Taylor-stabilized neural FS).

$$\begin{aligned} \text{FS}_{\text{nn}}(\alpha, h, m) \\ &= \frac{c + \phi h \cos^2 \alpha}{\gamma h \sin \alpha \cos \alpha + w_m m + \varepsilon}. \end{aligned} \quad (7)$$

Failure energy and the pixel gate are

$$\begin{aligned} \text{FE} &= \text{ReLU}(1 - \text{FS}_{\text{nn}}), \\ g &= \sigma(\psi - \text{FS}_{\text{nn}}). \end{aligned} \quad (8)$$

Equation (7) is deliberately *not* identical to (1): pore pressure has been moved from a subtractive numerator into an additive denominator by Proposition 1, and $(\tan \phi', \gamma_{\text{sat}}, \gamma_w)$ have been reparameterized into bounded learnable positives (Definition 2). This is the form that replaces generic affine neurons in mechanistic cells.

Definition 2 (Bounded material reparametrization). *For each material logit $w \in \{w_c, w_\phi, w_\gamma, w_m\}$,*

$$\text{positive_scale}(w) = \exp(\text{clip}(w; -8, 8)), \quad (9)$$

so $c, \phi, \gamma, w_m \in [e^{-8}, e^8]$. Clipping before the exponential prevents overflow while preserving positivity and smooth gradients almost everywhere.

3.4 Pixel and latent mechanistic cells

PixelMechanisticCell. At (or near) full spatial resolution, the DEM-derived angle α is a known geometric constant per pixel. Hardcoding $\sin \alpha \cos \alpha$ and $\cos^2 \alpha$ into the gate is therefore legitimate: the network need not rediscover the trigonometric landscape of gravity. The cell applies a 1×1 convolution and multiplies by the FS gate:

Definition 3 (PixelMechanisticCell).

$$\begin{aligned} & \text{PixelMechanisticCell}(x, \alpha, h, m) \\ &= \text{Conv}_{1 \times 1}(x) \\ & \odot \sigma(\psi - \text{FS}_{\text{nn}}(\alpha, h, m)). \end{aligned} \quad (10)$$

When (α, h, m) and x differ in spatial size, the proxies are bilinearly resized to match x . If mechanistic gating is disabled (ablation), the cell returns the convolution alone. Conceptually, low FS_{nn} (unstable slopes) opens the gate and admits features; high FS_{nn} (stable slopes) attenuates them.

LatentMechanisticCell. Deeper feature maps do not live in a geodetic coordinate frame. Applying $\sin \alpha$ or $\cos \alpha$ to abstract channels would force angular geometry onto tensors that no longer represent angles—exactly the pathology we seek to avoid. We therefore drop trigonometry and realize the same resisting/driving *ratio skeleton* with softplus projections:

Definition 4 (LatentMechanisticCell).

$$\begin{aligned} R &= \text{Softplus}(W_R * H), \\ D &= \text{Softplus}(W_D * H), \\ \text{LatentMechanisticCell}(H) \\ &= \text{Conv}_{3 \times 3} \left(\text{LeakyReLU} \left(\psi - \frac{R}{D + \varepsilon} \right) \right). \end{aligned} \quad (12)$$

Thus FS_{nn} and the latent ratio $R/(D + \varepsilon)$ belong to one continuum class: pixel cells keep DEM angles; latent cells keep only the continuum ratio. This dual realization is the operational meaning of Latent Continuum Equivalency at the cell level.

PhysicsEncoder. Each PhysicsEncoder stacks a pixel stem followed by successive latent cells with downsampling, then projects every level to width C :

$$\begin{aligned} f_0 &= \text{Stem}(\text{PixelMech}(x, \alpha, h, m)), \\ f_{\ell+1} &= \text{LatentMech}(\text{Down}(f_\ell)), \quad \ell = 0, \dots, 3, \\ P^L &= \text{Proj}_L(f_L), \quad L = 0, \dots, 4. \end{aligned} \quad (13)$$

At 256×256 input the physics pyramid lives on $\{256, 128, 64, 32, 16\}$. Two independent PhysicsEncoders run on RGB and DEM with their own proxies, yielding $\{P_{\text{rgb}}^L\}$ and $\{P_{\text{dem}}^L\}$.

3.5 Physics proxies and foundation inlet

Proxy extraction. Dataset channels are min–max normalized to $[0, 1]$. On Landslide4Sense, $s = x_2[:, 1]$, $d = x_2[:, 2]$, $v = x_2[:, 0]$. On Bijie (replicated DEM),

$$\begin{aligned} d &= x_2[:, 0], \\ s &= \text{Norm}(\|\nabla d\|), \\ v &= \text{Norm} \left(\frac{G-R}{G+R+\varepsilon} \right), \end{aligned} \quad (14)$$

where ∇d is a Sobel gradient and G, R are green/red bands of x_1 .

PhysicsProxyMapper. Separate RGB-path and DEM-path mappers (identical architecture, independent weights) produce

$$\begin{aligned} \alpha &= \frac{\pi}{2} \sigma(w_\alpha s + b_\alpha), \\ h &= \text{Softplus}(w_h d + b_h) + \varepsilon, \\ m &= \sigma(w_m v + b_m). \end{aligned} \quad (15)$$

Softplus keeps $h > 0$; sigmoid keeps $m \in (0, 1)$; the $\pi/2$ scale maps the slope logit into a valid angle range.

Foundation stack. A six-band observed stack feeds Prithvi-EO with LoRA [22, 23]: on Landslide4Sense, $[R, G, B, \text{NDVI}, \text{slope}, \text{DEM}]$; on Bijie, $[R, G, B, \text{DEM}, \text{Sobel}(d), \text{vegIndex}]$. Intermediate transformer blocks are projected to width C and resized onto a five-level pyramid $\{T_{\text{fm}}^L\}$.

3.6 Stage A–B: Pentastream encoding and Complementary Modality Bridge

Stage A runs five encoders in parallel:

1. E_{rgb} : EfficientNet [25] on RGB, pyramid $\{E_{\text{rgb}}^i\}_{i=0}^4$;
2. P_{rgb} : PhysicsEncoder on RGB with (α_r, h_r, m_r) ;
3. E_{dem} : EfficientNet on DEM, pyramid $\{E_{\text{dem}}^i\}$;
4. P_{dem} : PhysicsEncoder on DEM with (α_d, h_d, m_d) ;
5. T_{fm} : Prithvi+LoRA foundation pyramid.

EfficientNet-B4 at 256×256 yields spatial grids $\{128, 64, 32, 16, 8\}$ with growing channel counts. Physics and Prithvi pyramids use different native grids; they are bilinearly aligned to each EfficientNet level via a fixed legacy index map $\text{LEGACY} = (1, 2, 3, 4, 4)$, so EffNet index i consumes physics level $\text{LEGACY}[i]$. CNN texture is *complemented*, not replaced: CMB increases trust in physics only when CNN and physics co-activate (resonate).

Definition 5 (Complementary Modality Bridge). *Let $\tilde{E} = \text{GN}(\text{ReLU}(\text{Conv}_{1 \times 1}(E)))$ project the CNN tensor to width C and $\tilde{P} = \text{GN}(P)$ normalize the aligned physics tensor. Then*

$$g = \sigma(\text{Conv}_{1 \times 1}(\tilde{E} \odot \tilde{P})), \quad (16)$$

$$B_{\text{mix}} = \text{Conv}_{1 \times 1}([\tilde{E} \parallel \tilde{P}]), \quad (17)$$

$$H = g \odot \tilde{P} + (1 - g) \odot \tilde{E} + B_{\text{mix}}. \quad (18)$$

The gate g opens toward physics when element-wise agreement $\tilde{E} \odot \tilde{P}$ is strong; otherwise appearance features dominate. The residual mix B_{mix} preserves a shared channel pathway so neither stream is hard-gated to zero. Applying CMB at every EffNet index produces hybrid pyramids $\{H_{\text{rgb}}^i\}$ and $\{H_{\text{dem}}^i\}$. Foundation maps are likewise aligned to E_{rgb}^i , yielding $\{\hat{T}_{\text{fm}}^i\}$. Subsequent fusion consumes only these hybrids and foundation maps—not raw EffNet features. Algorithm 1 states the exact operator used in code.

Require: CNN features E , physics P

1. $P \leftarrow \text{match_spatial}(P, E)$
2. $\tilde{E} \leftarrow \text{GN}(\text{ReLU}(\text{Conv}_{1 \times 1}(E)))$
3. $\tilde{P} \leftarrow \text{GN}(P)$
4. $g \leftarrow \sigma(\text{Conv}_{1 \times 1}(\tilde{E} \odot \tilde{P}))$
5. $B \leftarrow \text{Conv}_{1 \times 1}([\tilde{E}; \tilde{P}])$
6. **return** $g \odot \tilde{P} + (1 - g) \odot \tilde{E} + B$

Algorithm 1: Complementary Modality Bridge (CMB)

3.7 Stage C: MAO-GeoEGCA (bottleneck fusion)

Stage C cross-references local geomechanics with planetary-scale context. At neck and bottleneck levels $i \in \{2, 3\}$ (spatial grids 32×32 and 16×16 for B4@256), Manifold-Aligned Orthogonal Geo-EGCA (**MAO** = manifold-aligned orthogonal; **GeoEGCA** = Geo-Equilibrium Gated Cross-Attention) treats the concatenated RGB/DEM hybrids as a *physics anchor* and Prithvi as environmental context. Restricting MAO to these two levels leaves four $\times 2$ upsampling stages in the decoder, restoring full 256×256 logits; applying MAO at 8×8 would under-resolve the decode chain.

Physical Equilibrium Gate. The gate G is a *pixel-wise* spatial mask at the bottleneck resolution (shape $B \times 1 \times H \times W$), not a single patch scalar. After Prithvi tokens are resized to standard 2D maps, G is formed by element-wise resonance $\hat{T} \odot X_p$, a 3×3 depthwise-separable neighbourhood filter, and a sigmoid. Low G suppresses planetary context where local physics indicate equilibrium; high G admits foundation context where instability cues are strong. The Value tensors V carry the foundation *content* that attention weights redistribute; without V , attention would score relevance but transport no payload.

Definition 6 (MAO-GeoEGCA). *Form the physics anchor and equilibrium gate*

$$\begin{aligned} X_p &= \text{Conv}_{1 \times 1}([\hat{H}_{\text{rgb}} \parallel \hat{H}_{\text{dem}}]), \\ G &= \sigma(\text{DWSepConv}_{3 \times 3}(\hat{T} \odot X_p)). \end{aligned} \quad (19)$$

Project queries from the anchor and keys/values from the foundation stream,

$$Q = W_q X_p, \quad K = W_k \hat{T}, \quad V = W_v \hat{T}. \quad (20)$$

Manifold alignment normalizes queries and modulates keys,

$$\hat{Q} = \frac{Q}{\|Q\|_2 + \varepsilon}, \quad K' = K \odot \hat{Q}, \quad (21)$$

so attention is biased toward directions consistent with the physics anchor. Multi-head attention (four heads, $d_h = C/4$) and gated residual readout give

$$\text{Attn} = \text{softmax}\left(\frac{QK'^\top}{\sqrt{d_h}}\right)V, \quad (22)$$

$$F = \text{Conv}_{3 \times 3}(\text{reshape}(\text{Attn}) \odot G) + X_p. \quad (23)$$

When local geomechanics indicate low stress, $G \rightarrow 0$ and foundation context is suppressed; when instability cues are strong, $G \rightarrow 1$ and Prithvi tokens guide the fused map. Outputs F^2 (neck, $B \times 64 \times 32 \times 32$) and F^3 (bottleneck, $B \times 64 \times 16 \times 16$) seed the dual decoder. Algorithm 2 matches the implemented fusion kernel.

Require: $\hat{T}, \hat{H}_{\text{rgb}}, \hat{H}_{\text{dem}}$

1. $X_p \leftarrow \text{Conv}_{1 \times 1}([\hat{H}_{\text{rgb}}; \hat{H}_{\text{dem}}])$
2. $G \leftarrow \sigma(\text{DWSepConv}_{3 \times 3}(\hat{T} \odot X_p))$
3. $Q \leftarrow W_q X_p; \quad K \leftarrow W_k \hat{T}; \quad V \leftarrow W_v \hat{T}$
4. $\hat{Q} \leftarrow \text{normalize}(Q); \quad K' \leftarrow K \odot \hat{Q}$
5. $\text{Ctx} \leftarrow \text{MHAttn}(Q, K', V)$
6. **return** $\text{Conv}_{3 \times 3}(\text{reshape}(\text{Ctx}) \odot G) + X_p$

Algorithm 2: MAO-GeoEGCA (neck / bottleneck)

3.8 Stage C (continued): TTEB skip lattice

For each fusion level $L \in \{0, 1, 2, 3\}$, the Tri-Temporal Tri-Stream Bridge builds a skip tensor from a lattice over RGB, DEM, and foundation streams at neighbouring scales $\{L - 1, L, L + 1\}$ (boundary-clamped). Levels are resized to the present spatial size. The present-phase tokens are mixed as

$$A = \text{Conv}_{1 \times 1}([\text{pres}_{\text{rgb}} \parallel \text{pres}_{\text{dem}} \parallel \text{pres}_{\text{fm}}]). \quad (24)$$

Eight off-diagonal lattice nodes (previous/next across streams and cross terms) are concatenated and mixed into a context tensor H_{ctx} . A latent FS residual—the same ratio skeleton as Definition 4—defines a stability map

$$\delta = \left| \psi - \frac{\text{Softplus}(W_R A)}{\text{Softplus}(W_D A) + \varepsilon} \right|. \quad (25)$$

Chunked spatial attention uses queries from A and keys/values from H_{ctx} , with an optional temporal router $\omega = \text{softmax}(W_\tau \delta)$ that emphasizes the present phase. The skip is

$$S^L = A + \text{Conv}_{3 \times 3}(\text{AttnOut} \odot \sigma(\delta)). \quad (26)$$

Thus TTEB does not invent a new inductive bias: its stability map δ is LCE applied to multi-scale, multi-stream skips. The resulting list $\{S^0, S^1, S^2, S^3\}$ lives on $\{128^2, 64^2, 32^2, 16^2\}$. Algorithm 3 states the per-level skip construction.

Require: hybrid/FM pyramids; level L

1. gather $\{L-1, L, L+1\}$ per stream; align spatially
2. $A \leftarrow \text{Conv}_{1 \times 1}([\text{pres}_{\text{rgb}}; \text{pres}_{\text{dem}}; \text{pres}_{\text{fm}}])$
3. $H_{\text{ctx}} \leftarrow \text{mix}(8 \text{ lattice nodes})$
4. $\delta \leftarrow |\psi - R(A)/(D(A) + \varepsilon)|$
5. $\text{AttnOut} \leftarrow \text{SpatialAttn}(A, H_{\text{ctx}}); \text{ scale by } \omega(\delta)$
6. **return** $A + \text{Conv}_{3 \times 3}(\text{AttnOut} \odot \sigma(\delta))$

Algorithm 3: TTEB at level L

3.9 Stage D: Dual Physics Decoder, PGDI, and MPEF

Stage D deliberately *re-splits* the fused Stage C knowledge into two competitive decode paths so spectral and topographic evidence can double-check the mask under commensurable continuum operators.

Shared vs. path-specific inputs. Both paths share $\{F^3, F^2, S^L\}$ from MAO/TTEB. Path A injects EfficientNet RGB residual $a_5 = \text{Proj}(E_{\text{rgb}}^4)$ with proxies (α_r, h_r, m_r) ; Path B injects DEM residual b_5 with (α_d, h_d, m_d) . The paths do not share hidden states during upsampling; they cooperate only at MPEF.

Dual paths. Both paths apply the same decode grammar—latent mechanistic cells, PGDI, and a final pixel mechanistic head—so evidence is decoded under one continuum class rather than a single shared statistical decoder.

PGDI. At each upsample stage, Physics-Gated Decoder Injection injects a TTEB skip through a learned gate on a latent mechanistic cell:

Definition 7 (PGDI). With decoder state D and aligned skip $S' = \text{align}(S, D)$,

$$\beta = \sigma(\text{Conv}_{1 \times 1}([D \parallel S'])), \quad (27)$$

$$D' = D + \beta \odot \text{LatentMechanisticCell}(S').$$

The gate β is content-dependent: skips that disagree with the current decode state are attenuated; skips that carry complementary continuum structure are admitted. Concretely, the decode chain widens $16^2 \rightarrow 32^2 \rightarrow 64^2 \rightarrow 128^2 \rightarrow 256^2$. At 32^2 , the upsampled state is first refined by a Latent Mechanistic Cell and then residual-injected with the MAO neck F^2 ; immediately afterward, PGDI admits the aligned TTEB skip. After four $\times 2$ stages and a pixel FS head, each path emits deeply supervised logits $(O_{\text{main}}, O_{\text{aux2}}, O_{\text{aux3}})$ of shapes $B \times 1 \times 256 \times 256$, $B \times 1 \times 64 \times 64$, and $B \times 1 \times 32 \times 32$. Algorithm 4 states the skip injection.

Require: decoder state D , skip S

1. $S' \leftarrow \text{match_spatial}(S, D)$
2. $\beta \leftarrow \sigma(\text{Conv}_{1 \times 1}([D; S']))$
3. **return** $D + \beta \odot \text{LatentMechanisticCell}(S')$

Algorithm 4: Physics-Gated Decoder Injection (PGDI)

MPEF. Dual-path logits are merged by Mechanistic Path Equilibrium Fusion—softmax failure energy of the same FS_{nn} derived in Sec. 3.3, not a DiGATE-style scalar GateFuse. Failure energy is computed from image-derived physics proxies (α, h, m) , bilinearly resized to each head resolution; it does *not* replace the decoder hidden state, but supplies the routing currency that adjudicates between candidate masks O_A and O_B :

Definition 8 (MPEF). For path logits O_A, O_B and proxies (α_p, h_p, m_p) ,

$$\text{FE}_p = \text{ReLU}(1 - \text{FS}_{\text{nn}}(\alpha_p, h_p, m_p)), \quad (28)$$

$$[w_A, w_B] = \text{softmax}([\text{FE}_A, \text{FE}_B]), \quad (29)$$

$$O = w_A O_A + w_B O_B + \text{Conv}_{1 \times 1}([O_A \parallel O_B]), \quad (30)$$

$$\mathcal{R}_{\text{MPEF}} = -\mathbb{E}[w_A \log(w_A + \varepsilon) + w_B \log(w_B + \varepsilon)]. \quad (31)$$

Independent MPEF modules fuse main, aux2, and aux3 heads (three masks during training; inference uses the main head). The entropy regularizer discourages indecisive averaging while preserving the continuum interpretation of w . Algorithm 5 states the merge; Algorithm 6 summarizes the full forward pass.

Require: O_A, O_B ; path proxies (α_p, h_p, m_p)

1. align proxies to logit resolution
2. $\text{FE}_p \leftarrow \text{ReLU}(1 - \text{FS}_{\text{nn}}(\alpha_p, h_p, m_p))$
3. $[w_A, w_B] \leftarrow \text{softmax}([\text{FE}_A, \text{FE}_B])$
4. $O \leftarrow w_A O_A + w_B O_B + \text{Conv}_{1 \times 1}([O_A; O_B])$
5. $R \leftarrow \text{mean}(-(w \log(w + \varepsilon)))$
6. **return** O, R

Algorithm 5: Mechanistic Path Equilibrium Fusion (MPEF)

Require: streams x_1 (RGB), x_2 (topo/aux)

1. $(s, d, v) \leftarrow \text{physics_proxies}(x_1, x_2)$; map (α, h, m)
2. encode $E_{\text{rgb}}, E_{\text{dem}}, P_{\text{rgb}}, P_{\text{dem}}, T_{\text{fm}}$
3. **for** $i = 0.4$: CMB hybrids H^i ; align $\hat{T}^i$
4. **for** $L \in \{0.3\}$: $S^L \leftarrow \text{TTEB}(\dots; L)$
5. $F^2, F^3 \leftarrow \text{MAO}$ at levels 2, 3
6. dual PhysDec $\rightarrow (O_A, O_B)$ at three heads
7. MPEF merge each head; **return** logits+routing reg

Algorithm 6: PS-GPLNet forward

3.10 Training objective

We use deeply supervised Tversky loss [26] on sigmoid logits,

$$\mathcal{T}(p, t) = \frac{TP + s}{TP + \alpha_{\text{tv}} FP + \beta_{\text{tv}} FN + s}, \quad (32)$$

$$\mathcal{L}_{\text{TV}} = 1 - \mathcal{T},$$

with defaults $\alpha_{\text{tv}} = 0.3$, $\beta_{\text{tv}} = 0.7$ (recall emphasis; swapped on Bijie where precision is preferred). Heads are bilinearly upsampled to input resolution before loss evaluation. The total objective is

$$\mathcal{L} = \sum_{h \in \{\text{main}, \text{aux2}, \text{aux3}\}} w_h \mathcal{L}_{\text{TV}}^{(h)} + \lambda_{\text{reg}} \sum_h \mathcal{R}_{\text{MPEF}}^{(h)}, \quad (33)$$

with $(w_{\text{main}}, w_{\text{aux2}}, w_{\text{aux3}}) = (1.0, 0.6, 0.4)$ and $\lambda_{\text{reg}} = 10^{-3}$. Training uses Adam ($3 \cdot 10^{-4}$, weight decay 10^{-4}), up to 100 epochs, frozen pretrained EfficientNet towers when available, and frozen Prithvi with LoRA adapters.

4 Theoretical Analysis

We now state the claims that organize Latent Continuum Equivalency. The continuum class $\mathcal{C}_{\text{FS}}$ comprises Definition 1 at pixels and Definition 4 in latent space—the same resisting/driving ratio after Proposition 1. LCE is what makes failure energy a shared currency from inlet gating through MPEF routing.

Definition 9 (Latent Continuum Equivalency). *A layered multimodal network satisfies LCE if every stage ℓ applies an operator $T_\ell \in \mathcal{C}_{\text{FS}}$ (identical algebraic FS skeleton, possibly different parameters and resolutions), so failure energies remain commensurable across depths and paths.*

Lemma 1 (Compactness of materials). *Under Definition 2, $c, \phi, \gamma, w_m \in [e^{-8}, e^8]$.*

Proof. $\text{clip}(w; -8, 8)$ lands in $[-8, 8]$; $\exp$ is continuous and strictly increasing, so the image is exactly $[e^{-8}, e^8]$. $\square$

Lemma 2 (Local Lipschitz continuity of FS_{nn}). *On any compact training domain $\alpha \in [\alpha_{\min}, \alpha_{\max}] \subset (0, \pi/2)$, $h \geq h_{\min} > 0$, $m \in [0, 1]$, the denominator of (7) is at least $\varepsilon > 0$, hence FS_{nn} is C^∞ and Lipschitz on that set.*

Proof. The map is a ratio of smooth functions with non-vanishing denominator. Continuous functions on compact metric spaces are Lipschitz [14, 19]. $\square$

Lemma 3 (FS gate Lipschitz). *$x \mapsto \sigma(\psi - \text{FS}_{\text{nn}}(x))$ is Lipschitz on the same domain with constant at most $\frac{1}{4}\text{Lip}(\text{FS}_{\text{nn}})$, since $\text{Lip}(\sigma) \leq 1/4$.*

Proof. Composition of Lipschitz maps; $\sup |\sigma'| = 1/4$. $\square$

Proposition 2 (Non-explosion of the pixel cell). *Composing a Lipschitz 1×1 map f with Lemma 3 yields a locally Lipschitz mechanistic cell. Unbounded $\exp(w)$ without clipping is excluded by Lemma 1. Together, Proposition 1 and Lemmas 1–3 guarantee that the neural FS operator remains a well-posed drop-in replacement for affine neurons under standard training domains.*

Theorem 1 (LCE enables FE routing). *If every path and stage applies $\mathcal{C}_{\text{FS}}$ under Lemma 1, then path-wise FE_A, FE_B are evaluations of the same continuum currency and are admissible inputs to Definition 8. Without LCE, dual-path weights lack a shared physical unit.*

Proof. Commensurability follows from identical FS algebra (Proposition 1 $\rightarrow$ Definitions 1,4) and shared positivity (Lemma 1). Arbitrary learned gates are not evaluations of $\mathcal{C}_{\text{FS}}$ and therefore are not FE-comparable. $\square$

Theorem 2 (MPEF energy identity). *On the simplex $\Delta = \{w_A + w_B = 1, w \geq 0\}$, $\mathbf{w}^* = \text{softmax}(\mathbf{FE})$ uniquely maximizes $\mathbf{w}^\top \mathbf{FE} + \mathcal{H}(\mathbf{w})$.*

Proof. This is the standard entropic projection / soft-argmax identity on a two-point simplex [14]. Strict concavity of Shannon entropy $\mathcal{H}$ yields uniqueness. $\square$

Corollary 1 (Contrast to GateFuse). *Scalar learned fusion gates that are not $\text{softmax}(\mathbf{FE})$ are not solutions of Theorem 2; MPEF is an equilibrium of continuum failure energy, not an unconstrained gate [12, 9].*




Figure 2: Bounded material map $\exp(\text{clip}(w))$ versus unbounded $\exp(w)$.




Figure 3: Infinite-slope FS landscape over (α, h) .

Remark 1 (Empirical probes). *On GPU we verified Lemma 1 numerically ($c \in [e^{-8}, e^8]$), Theorem 2 against 5000 random simplex competitors (dominance rate 1.0), and local FS-gate sensitivity (median $\|\Delta \text{gate}\| / \|\Delta \alpha\| \approx 0.28$, $p_{95} \approx 0.68$). See Figures 2–4.*

Remark 2 (On Universal Approximation). *We do not claim a Universal Approximation Theorem for continuum-in-neuron operators. UAT would establish existence of some network approximating continuous maps; our guarantees are instead local well-posedness and routing identities (Lemmas 1–3, Theorems 1–2) that make LCE/MPEF drop-in operators under standard training domains. This is the appropriate theory depth for TPAMI-style multimodal dense prediction.*

5 Experiments

5.1 Stress-test datasets and protocol

We evaluate on Bijie [15] and Landslide4Sense [16] as anisotropic multimodal dense-prediction stress tests. Inputs




Figure 4: MPEF path weights as functions of relative failure energy.




Figure 5: Empirical local sensitivity of the FS gate (GPU-0 Monte Carlo).




Figure 6: Schematic of Latent Continuum Equivalency: shared FS currency from inlet to MPEF.

are 256×256 . Adam ($3 \cdot 10^{-4}$, weight decay 10^{-4}), Tversky loss, 100 epochs, threshold 0.5. EfficientNet backbones frozen when pretrained; Prithvi frozen with LoRA. Metrics: Acc/Prec/Rec/F1/IoU on the main head. Final numbers are best validation epochs from *outputs_final_bijie* (epoch 92) and *outputs_final_l4s* (epoch 70).




Figure 7: Validated gains of PS-GPLNet over DeepLabV3+ and dual_stream_gated.




Figure 8: IoU ranking at best epoch on both stress-test datasets.




Figure 9: Bijie grouped best-epoch metrics including PS-GPLNet.

5.2 Baselines

LinkNet, DEP-UNet, GMNet, RMAU-Net, TransUNet, EMR-HRNet, ShapeFormer, Dual-Stream UNet, U-Net, DeepLabV3+, and dual_stream_gated under a unified harness. DiGATE-style dual-stream gated models are baselines only [12].

5.3 Quantitative results




Table 2 consolidates both benchmarks. On Bijie, PS-GPLNet reaches $\text{F1} = 0.933$ and $\text{IoU} = 0.907$, above DeepLabV3+ (0.927/0.899) and dual_stream_gated (0.897/0.869). On Landslide4Sense, $\text{F1} = 0.736$ and $\text{IoU} = 0.660$, above DeepLabV3+ (0.710/0.636). Gains versus DeepLabV3+ and dual-gated baselines are summarized in Figure 7 (Bijie $\Delta \text{F1} + 0.006/\Delta \text{IoU} + 0.009$; L4S $\Delta \text{F1} + 0.026/\Delta \text{IoU} + 0.024$ vs. DeepLab).

5.4 Qualitative stress-test results

Figures 13–14 show representative patches. Bijie polygons are large; heatmaps concentrate on scarps and deposits while

Table 2: Best-epoch validation on Bijie and Landslide4Sense (unified protocol).

Model	Bijie					Landslide4Sense				
	Acc	Prec	Rec	F1	IoU	Acc	Prec	Rec	F1	IoU
LinkNet	—	—	—	—	—	0.973	0.487	0.735	0.431	0.359
DEP-UNet	0.982	0.850	0.934	0.831	0.799	0.977	0.669	0.770	0.603	0.535
RMAU-Net	0.981	0.859	0.924	0.830	0.798	0.977	0.691	0.756	0.618	0.548
ShapeFormer	0.984	0.892	0.899	0.840	0.810	0.975	0.675	0.802	0.632	0.560
TransUNet	0.983	0.889	0.908	0.845	0.815	0.976	0.714	0.750	0.619	0.551
EMR-HRNet	0.983	0.903	0.905	0.846	0.818	0.975	0.674	0.784	0.627	0.557
GMNet	0.983	0.900	0.909	0.855	0.825	0.978	0.669	0.779	0.616	0.545
U-Net	0.987	0.891	0.933	0.873	0.842	0.984	0.795	0.719	0.680	0.611
Dual-Stream UNet	0.986	0.902	0.928	0.875	0.843	0.986	0.713	0.779	0.656	0.581
dual_stream_gated	0.988	0.928	0.929	0.897	0.869	0.984	0.822	0.673	0.646	0.573
DeepLabV3+	0.989	0.929	0.966	0.927	0.899	0.984	0.789	0.749	0.710	0.636
PS-GPLNet	0.991	0.956	0.946	0.933	0.907	0.986	0.802	0.788	0.736	0.660

**Figure 10:** Landslide4Sense grouped best-epoch metrics including PS-GPLNet.**Figure 11:** Bijie final-run training dynamics (best F1 at epoch 92).**Figure 12:** Landslide4Sense final-run training dynamics (best F1 at epoch 70).

suppressing homogeneous steep slopes. L4S scars are smaller and cross-region; precision–recall trade-offs match Table 2.

5.5 Theory–experiment alignment

Material compactness, MPEF uniqueness, and FS-gate sensitivity (Remark after Theorem 2) were verified on an RTX 4000 Ada (20 GB) without retraining. Training curves (Figures 11–12) show stable validation F1/IoU on the final


PS-GPLNet logs.

5.6 Component ablation study

Following the diagnostic organization used in strong multi-modal ablations (Difficult vs. Easy stress tests; component matrix; mechanism swaps), we treat **Landslide4Sense as the hard regime** (cross-region scars, severe imbalance) and **Bijie as the easier regime** (larger polygons, more homogeneous acquisition). Absolute Full-model numbers in Table 3 are taken from the primary-server runs already reported in Table 2 (Bijie F1/IoU 0.933/0.907; L4S 0.736/0.660)—*not* from the compact ablation Full checkpoint. Component rows are matched-protocol removals under a compact low-VRAM schedule (EfficientNet-B0, $C=32$, LoRA $r \leq 4$, FP32, batch 2, 40 epochs, seed 42). Protocol-matched Δ and mechanism swaps appear in the bottom of Table 3 (compact Full reference: Bijie F1 0.804; L4S F1 0.693); comparative plots are in Figure 15.


What each removal tests. w/o FS-gate disables Taylor-stabilized FS gating (identity residual). w/o MPEF replaces instability routing by mean fusion. w/o CMB bypasses Complementary Modality Bridge. w/o MAO replaces GeoEGCA by a plain 1×1 concat mix. w/o TTEB replaces the skip lattice by bilinear identity skips. w/o Prithvi drops the foundation stream. Path-A only keeps a single RGB-proxy decode path.

Hard-set verdict (L4S). The dominant failure mode under matched removal is dropping Prithvi ($\Delta F1 -0.163$, $\Delta IoU -0.159$): planetary context is indispensable when scars are small and acquisition is heterogeneous. Mean fusion without MPEF is a mild regression (-0.005 F1). Dual-path competition is not the primary L4S bottleneck under the compact budget (Path-A only is slightly above compact Full), consistent with foundation context dominating local path rivalry on the hard set.



(a) Bijie Landslide Dataset

Figure 13: Qualitative Bijie samples: RGB, ground truth, prediction, heatmap.



(c) Landslide4Sense Dataset

Figure 14: Qualitative Landslide4Sense samples.

Table 3: Ablation study on Difficult (Landslide4Sense) vs. Easy (Bijie). **Top:** Acc/Prec/Rec/F1/IoU; Full = primary-server PS-GPLNet (Table 2); other rows = matched compact protocol. **Bottom-left:** $\Delta F1/\Delta IoU$ vs. compact Full (Bijie 0.804/0.773; L4S 0.693/0.617). **Bottom-right:** mechanism-swap F1 under the same compact protocol.

Variant	Difficult: Landslide4Sense					Easy: Bijie				
	Acc	Prec	Rec	F1	IoU	Acc	Prec	Rec	F1	IoU
Full (primary)	0.986	0.802	0.788	0.736	0.660	0.991	0.956	0.946	0.933	0.907
w/o Prithvi	0.981	0.616	0.750	0.530	0.458	0.975	0.961	0.754	0.728	0.717
Path-A only	0.985	0.719	0.842	0.704	0.627	0.982	0.834	0.867	0.753	0.716
w/o CMB	0.985	0.755	0.809	0.715	0.638	0.987	0.801	0.960	0.795	0.767
w/o MPEF	0.983	0.691	0.853	0.689	0.611	0.987	0.860	0.918	0.812	0.782
w/o MAO	0.986	0.756	0.805	0.717	0.641	0.987	0.933	0.898	0.870	0.835
w/o TTEB	0.984	0.721	0.835	0.706	0.628	0.987	0.965	0.891	0.889	0.863
w/o FS-gate	0.984	0.736	0.807	0.703	0.626	0.988	0.948	0.916	0.896	0.869

Removed	Difficult (L4S)		Easy (Bijie)		Mechanism setting (F1)		
	$\Delta F1$	ΔIoU	$\Delta F1$	ΔIoU	Setting	L4S	Bijie
w/o Prithvi	-0.163	-0.159	-0.076	-0.057	Full (compact)	0.693	0.804
Path-A only	+0.011	+0.010	-0.052	-0.057	w/o FS-gate	0.703	0.896
w/o CMB	+0.022	+0.022	-0.009	-0.007	w/o MPEF	0.689	0.812
w/o MPEF	-0.005	-0.006	+0.007	+0.009	Path-A only	0.704	0.753
w/o MAO	+0.024	+0.025	+0.065	+0.062			
w/o TTEB	+0.013	+0.011	+0.085	+0.090			
w/o FS-gate	+0.010	+0.009	+0.092	+0.096			

Easy-set verdict (Bijie). Prithvi removal remains harmful (-0.076 F1), and Path-A only collapses dual-path competition (-0.052 F1), showing that topographic decode authority matters when polygons are large enough for DEM proxies to adjudicate. CMB removal is near-neutral under the short budget.

Positive Δ under compact protocol. Several removals (FS-gate, TTEB, MAO on Bijie; MAO/CMB on L4S) exceed compact Full under the 40-epoch low-VRAM schedule. We do *not* interpret this as “Full is worse”: primary Full (Table 2) still dominates every component absolute number. The pattern matches a known optimization interaction—strong continuum regularizers and multi-module fusion can underfit when capacity, LoRA rank, and epoch budget are jointly reduced—while the same operators recover and lead when trained under the primary protocol. Figure 15 summarizes F1/IoU, matched Δ , signed contributions, mechanism swaps, and the easy–hard ranking.




Figure 15: Comparative ablation panel. Dashed lines / star markers denote primary-server Full (Bijie 0.933, L4S 0.736).

5.7 Latent manifold, boundary fidelity, and perturbation robustness

Pixel IoU alone does not probe whether LCE structures the representation space or whether FS-gated decoding resists sensor noise. We therefore run three inference-only studies on the Bijie validation split (seed 42, 554 images) using the real GeoPhysicsLandslideNet forward API and a locally trained architecture-faithful B4 checkpoint (epoch 98, image-mean val F1 0.912; cluster final weights for the 0.933 table entry were not synced to this machine). Latents are captured by a forward hook on MAO `fuse3` ($C=64$), not 1-channel MPEF logits.

Latent manifold (t-SNE). Figure 16 projects 300 pooled `fuse3` vectors. Color by landslide presence ($> 5\%$ mask area) yields visible class structure (centroid separation 0.522); color by DEM channel std shows topography-aligned texture variation along the same embedding, consistent with continuum-aware routing rather than purely appearance clustering.

Boundary complexity. We report Boundary F-score ($\theta=2$ px) and symmetric Hausdorff distance on validation predictions (Figure 17). Over the full val set, mean BF-score is 0.821; restricting to landslide-present images ($n=154$) gives BF 0.394 and median Hausdorff 17.3 px, isolating anisotropic scar/deposit boundaries from trivial empty-empty zeros.

Perturbation decay. Gaussian noise $\sigma \in \{0, 0.05, 0.1, 0.2, 0.3\}$ is added to both RGB and DEM




Figure 16: t-SNE of MAO `fuse3` latents on Bijie val (hook $C=64$). Left: landslide presence; right: DEM texture proxy.




Figure 17: Boundary F-score and Hausdorff histograms on Bijie val predictions (PS-GPLNet).

streams. Figure 18 compares micro-averaged F1 decay against trained Bijie U-Net and DeepLabV3+ checkpoints from the same baseline harness. At $\sigma=0$, DeepLab starts highest (0.799) vs. PS-GPLNet micro-F1 0.759 and U-Net 0.736. Under strong noise ($\sigma=0.3$), baselines collapse to ≈ 0.49 – 0.50 while PS-GPLNet remains at 0.653, and at moderate $\sigma \in \{0.05, 0.1\}$ PS-GPLNet peaks near 0.827—evidence that FS/MPEF priors act as a structural anchor rather than a free statistical gate.




Figure 18: F1/IoU under RGB+DEM Gaussian noise. PS-GPLNet decays slower than U-Net and DeepLabV3+.

6 Discussion

The paper’s primary claim is methodological: LCE injects a continuum operator into multimodal dense predictors with bounded Jacobians, and MPEF turns dual-path fusion into an energy equilibrium. Landslide benchmarks are chosen because topographic–spectral fusion is anisotropic and imbalanced—a harder regime than RGB–D for testing continuum closure [4, 5]. Ablations (Sec. 5.6) sharpen that claim: on the Difficult set, foundation context is the largest matched drop; on the Easy set, dual-path competition and Prithvi both matter; positive compact-protocol Δ for some removals are disclosed as under-capacity interactions, while primary Full remains the absolute reference. Manifold/boundary/robustness studies further support the CV claim beyond application tables: latents organize by landslide and topography, boundaries are measurable with BF/Hausdorff, and physics-closed decoding resists input corruption better than strong statistical segmenters.

Do we need a Universal Approximation Theorem? No. For TPAMI / NeurIPS / ICCV-style multimodal dense prediction, UAT is neither necessary nor the right certificate: classical UAT concerns existence of approximators for continuous functions on compact sets, whereas our contribution is an *operator family* with compactness of materials (Lemma 1), local Lipschitz FS gates (Lemmas 2–3), and an MPEF energy identity (Theorem 2). Proving that continuum equations can “replace” neurons in the UAT sense would not strengthen acceptance relative to mechanism ablations, baseline comparisons, and robustness probes already provided. We likewise do not require training unrelated domains for the first submission if LCE+MPEF is framed as a reusable continuum-closure recipe and stress-tested on anisotropic multimodal segmentation; cross-domain transfer remains future work, not a gatekeeping theorem.

CV vs. application framing. The manuscript is positioned as physics-closed multimodal dense prediction (LCE/MPEF), with landslide inventories as stress tests—not as a susceptibility GIS product. Limitations remaining: local vs. clus-

ter checkpoint sync for the exact 0.933 weights, compact vs. primary training budgets in ablations, and the fact that FS is a reduced-order prior rather than a full 3D slope-stability PDE [24, 27].

7 Conclusion

We presented Latent Continuum Equivalency and Mechanistic Path Equilibrium Fusion as a physics-closed recipe for multimodal dense prediction, instantiated as PS-GPLNet and stress-tested on Bijie and Landslide4Sense. Theory guarantees compactness and local Lipschitz behavior of FS cells and identifies MPEF with entropic energy maximization—without invoking Universal Approximation as a prerequisite. Experiments show strong primary F1/IoU under a unified baseline harness, Difficult/Easy component ablations with protocol-matched Δ , plus latent-manifold structure, boundary metrics, and flatter noise-decay curves than U-Net/DeepLabV3+. Future work includes syncing final cluster checkpoints and transferring LCE cells beyond geohazard segmentation.

Acknowledgements

We thank CSIR-4PI and KIIT for research support, and the maintainers of Bijie, Landslide4Sense, Prithvi-EO, and open baseline implementations.

References

- [1] Olaf Ronneberger, Philipp Fischer, and Thomas Brox. U-net: Convolutional networks for biomedical image segmentation. *MICCAI*, 2015.
- [2] Liang-Chieh Chen, Yukun Zhu, George Papandreou, Florian Schroff, and Hartwig Adam. Encoder-decoder with atrous separable convolution for semantic image segmentation. In *ECCV*, 2018.
- [3] Enze Xie et al. Segformer: Simple and efficient design for semantic segmentation with transformers. 2021.
- [4] Wei Yin, Yifan Liu, and Chunhua Shen. Virtual normal: Enforcing geometric constraints for accurate and robust depth prediction. *IEEE Trans. Pattern Anal. Mach. Intell.*, 2021.
- [5] Yi Wang, Zhenghong Wang, Fan Zhang, et al. A gravity-informed spatiotemporal transformer for human activity intensity prediction. *IEEE Trans. Pattern Anal. Mach. Intell.*, 2025. arXiv:2506.13678.
- [6] Maziar Raissi, Paris Perdikaris, and George E Karniadakis. Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 2019.
- [7] George Em Karniadakis et al. Physics-informed machine learning. *Nature Reviews Physics*, 3:422–440, 2021.
- [8] James R. Clough, Nicholas Byrne, Ilkay Oksuz, et al. A topological loss function for deep-learning based image segmentation using persistent homology. *IEEE Trans. Pattern Anal. Mach. Intell.*, 2022.

- [9] Gargi Panda, Soumitra Kundu, Saumik Bhattacharya, and Aurobinda Routray. ℓ_0 -regularized sparse coding-based interpretable network for multi-modal image fusion. *IEEE Trans. Pattern Anal. Mach. Intell.*, 2025.
- [10] Ashok Dahal and Luigi Lombardo. Towards physics-informed neural networks for landslide prediction. *Engineering Geology*, 337:107852, 2024.
- [11] Songlin Liu, Luqi Wang, Wengang Zhang, et al. A physics-informed data-driven model for landslide susceptibility assessment in the three gorges reservoir area. *Geoscience Frontiers*, 2023.
- [12] Md Islam et al. Digate: Dual-stream gated attention network for landslide segmentation. *Remote Sensing*, 2024.
- [13] Alexandre Mercier et al. Bifusion-net: A multi-stream fusion network for landslide detection using high-resolution remote sensing imagery and dem. *Remote Sensing*, 14(3):612, 2022.
- [14] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [15] Shunping Ji, Dawen Yu, Chaoyong Shen, Weile Li, and Qiang Xu. Landslide detection from an open satellite imagery and digital elevation model dataset. *Landslides*, 2020.
- [16] Omid Ghorbanzadeh et al. Landslide4sense: Reference benchmark data and deep learning models for landslide detection. *IEEE Transactions on Geoscience and Remote Sensing*, 2022.
- [17] Jieneng Chen et al. Transunet: Transformers make strong encoders for medical image segmentation. *arXiv preprint arXiv:2102.04306*, 2021.
- [18] Omid Ghorbanzadeh et al. Evaluation of different machine learning methods and deep-learning convolutional neural networks for landslide detection. *Remote Sensing*, 2019.
- [19] Anonymous. Fast and flexible robustness certificates for semantic segmentation. *arXiv preprint arXiv:2512.06010*, 2025.
- [20] Diyar Altinses and Andreas Schwung. Layer-specific lipschitz modulation for fault-tolerant multimodal representation learning. *arXiv preprint arXiv:2603.25103*, 2026.
- [21] A. Maurya et al. Landslide susceptibility mapping using physics-guided machine learning. *Natural Hazards*, 2024.
- [22] Daniela Szwarcman et al. Prithvi-eo-2.0: A versatile multi-temporal foundation model for earth observation applications. *arXiv preprint*, 2024.
- [23] Neil Houlsby et al. Parameter-efficient transfer learning for nlp. 2019.
- [24] A W Skempton and F A DeLory. Stability of natural slopes in london clay. *Proc. 4th Int. Conf. Soil Mech.*, 1957.
- [25] Mingxing Tan and Quoc Le. Efficientnet: Rethinking model scaling for convolutional neural networks. 2019.
- [26] Seyed Sadegh Mohseni Salehi, Deniz Erdogmus, and Ali Gholipour. Tversky loss function for image segmentation using 3d fully convolutional deep networks. 2017.
- [27] Nathan M Newmark. Effects of earthquakes on dams and embankments. *Géotechnique*, 1965.